


# Ames Housing Data: Machine Learning Applications



---

Victoria Baker, Jan Forslow, Alice Lam, Josh Lee



# Overview

The data, and our approach



# The Data

---

- Feature counts:
  - 72 => 212 after processing
- Row counts:
  - 1460 test, 1459 train
- Features cover:
  - Size
  - Conditions rating
  - Utilities
  - Construction features
  - Community and neighborhood
- Data split for assessment:
  - Fit/Predict: mini\_train (80%), dev (20%)
  - Cross validation on mini\_train
  - Scoring: Root mean squared log error (RMSLE)



# Our Approach

Looking for the best model




Experimenting with other models





# Our Best Results (RMSLE)

RMSLE



**What have we  
learned  
from the process?**

*Why some work while some don't?*

**?**

1

# Key Learnings


From: EDA, Feature Processing



# Impact of feature processing

Based on Basic LinearRegression() Performance:

Fit mini\_train (N=1162)  
Pred dev (N = 291)






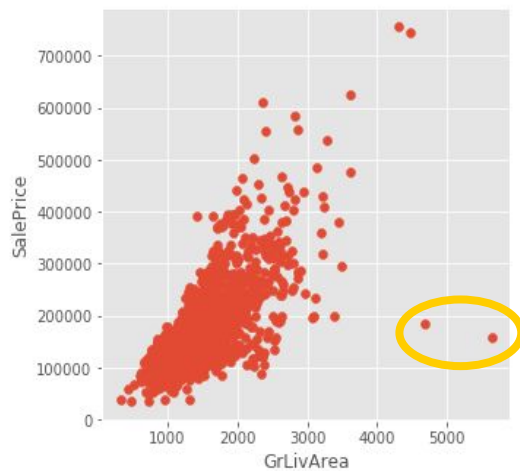


# Feature processing

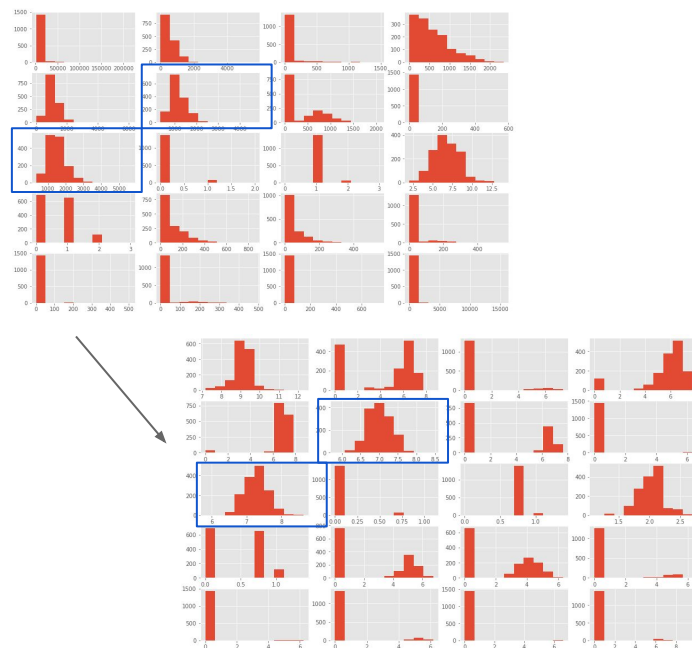
Taking Log of Y (SalePrice)



Outliers



Normalize skewness



---

2


# Key Learnings

From: Feature Engineering and GSCV

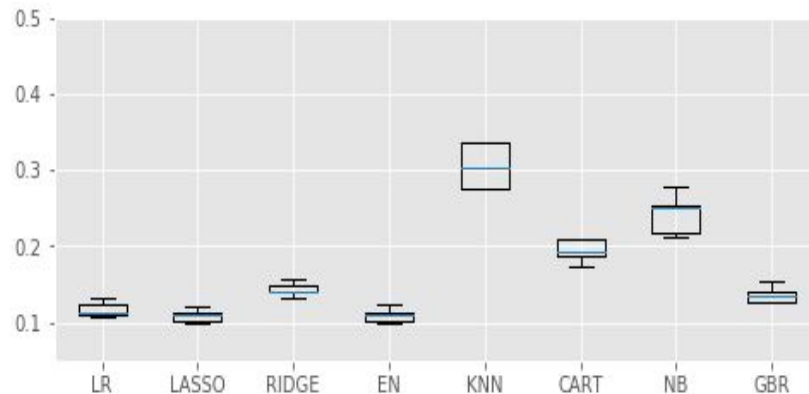


# GSCV: Comparing across models

Default Param



Best Param from GridsearchCV



**Best Param:**

**LR:** Normalize = False

**Lasso:** Normalize = True, Alpha = 0.0001

**Ridge:** Normalize = True, Alpha = 1.0

**EN (Elastic Net):** Fit\_Intercept= True, Normalize = True, Alpha = 0.0001, L1\_ratio = 0.9

**KNN:** N = 1

**CART (Decision Tree Regressor):** Max\_features = 10, max\_depth = 11, min\_samples\_split = 60


**NB:** alpha = 0.0001

**GBR (Gradient Boosting Regressor):** Learning\_rate = 0.04, n\_estimators = 100



# Feature Engineering

RMSLE on Pred(Dev\_Data) using different feature engineering methods



- No FE
- MinMaxScaler
- StandardScaler
- L1 Selection
- PCA only
- PCA + GMM
- GMM
- PCA + Kmeans
- K Means
- GMM grouped by SalesPrice
- PCA - size features
- Stack All

## Best Param:

**LR:** Normalize = False

**Lasso:** Normalize = True,  
Alpha = 0.0001

**Ridge:** Normalize = True,  
Alpha = 1.0

**EN:** Fit\_Intercept= True,  
Normalize = True,  
Alpha = 0.0001,  
L1\_ratio = 0.9


3

# Key Learnings

From: Neural Network



## NN on Regression Problem



4

# Key Learnings


From: Gradient Boosting & F Regression




# Helpful for Feature Learning

Feature Importance given by GBR

Correlation Matrix based on F-Regression, n=60



- GBR RMSLE of 0.12 at 500 trees, depth = 8, lr = 0.2 and 0.3 subsampling,
- Built-in function for feature importance
- Sub-selecting 60 features of 212 with f-regression to avoid multi-collinearity, still achieves 0.13 RMSLE





5

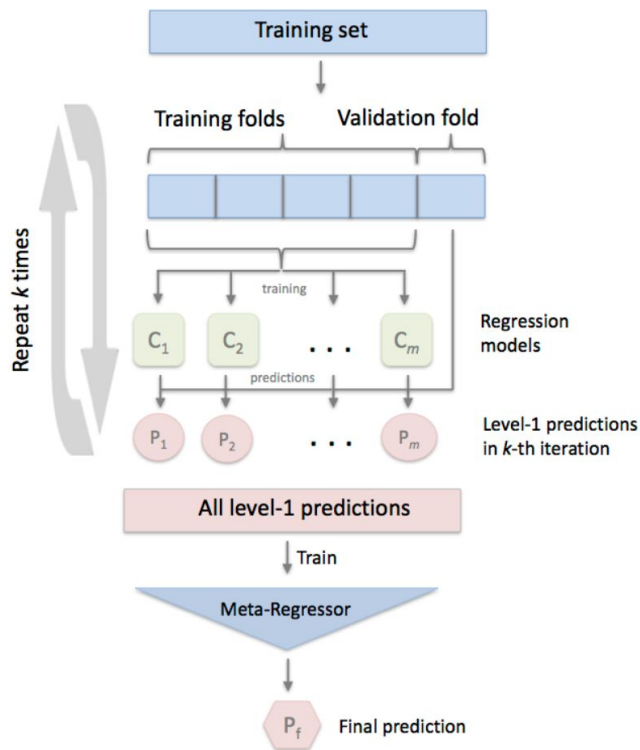
# Key Learnings

From: Ensembles



## Ensembles Hard to Find

- Tested three off-the-shelf ensemble modules
- Scikit-learn ensemble.VotingClassifier
  - Wraps models and averages the predictions
  - Not working for LR models
- mixtend.StackingRegressor
  - Individual regression models trained on complete training set; Meta-regressor LR fitted on the outputs
  - Combined LR, LASSO, EN, Ridge, CART and GBR
- mixtend.StackingCVRegressor
  - Using out-of-fold predictions to prepare the input data for the level-2 classifier; Avoids over-fitting
  - Achieved RMSLE of 0.11 on dev-set



6

# Key Learnings

From: One Drop Model & Decision Trees



# One Drop Model: What We Learned

## Key Points

1. Feature processing is very important
2. Use “system” level not “unit” level metrics

### RR Before FP

	0
count	212.000000
mean	0.418923
std	0.000114
min	0.418449
25%	0.418919
50%	0.418919
75%	0.418919
max	0.420471

### RR After FP

	0
count	2.1e+02
mean	2.4e-01
std	3.7e-02
min	1.4e-01
25%	2.4e-01
50%	2.5e-01
75%	2.5e-01
max	3.3e-01



# Decision Tree: Information Gain & f-regression

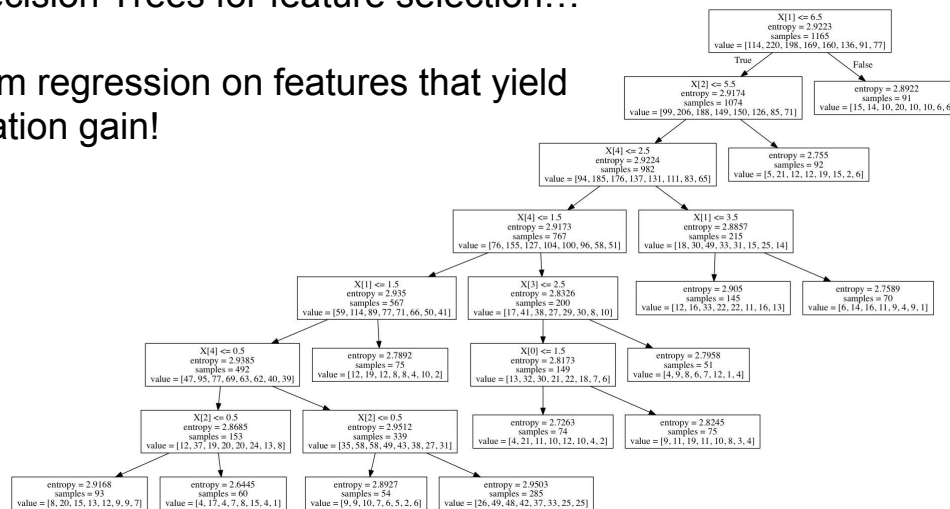
DT Similar to GBR  
Feature Importance


GBR	DT
OverallQual	OverallQual
GrLivArea	TotalBsmtSF
BsmtUnfSF	GrLivArea
TotalBsmtSF	GarageArea
1stFlrSF	GarageFinish
BsmtFinSF1	GarageType__Detchd

## Idea

Why not use Decision Trees for feature selection...

And then perform regression on features that yield greatest information gain!





# Conclusion

Ways to improve further



## Regression Diagnostics

---

- Problems meeting multicollinearity and homoscedasticity assumptions
- Difficult to combine/remove enough features to make an impact without losing potentially important data
- Feature engineering work should be informed by regression diagnostics, not limited by them



## Further Improvements

---

- Use regression diagnostics to optimize One Drop and f-regression models
- Run One Drop model with more features being dropped
- Try using Gradient Descent and see if it improves basic linear regression
- Design our own ensemble function that combines all model types
- Use Independent Component Analysis to decompose features with high mutual information
- Adding jitter to the features in training data to test stability of OLS, Lasso, Ridge, and Elastic Net models





# Thanks!

Any **questions** ?

Reference:


- Effects of multicollinearity:  
<http://blog.minitab.com/blog/adventures-in-statistics-2/what-are-the-effects-of-multicollinearity-and-when-can-i-ignore-them>

# Ames Housing Data: Machine Learning Applications



Victoria Baker, Jan Forslow, Alice Lam, Josh Lee


# **Additional Slides**

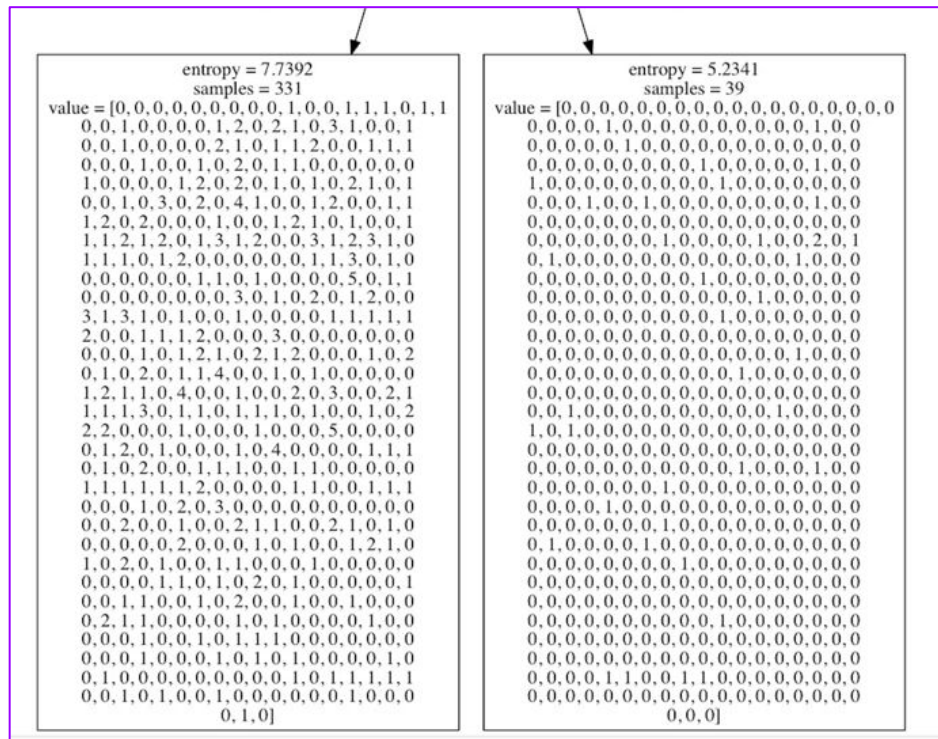








## Mutual Information








# Decision Trees: what happens when you use the default settings





# Feature Engineering

RMSLE on Pred(Dev\_Data) on 4 types of non-regression models



## Best Param:

**KNN:** N = 1

**CART:** Max\_features = 10,  
max\_depth = 11,  
min\_samples\_split = 60

**NB:** alpha = 0.0001

**GBR:** Learning\_rate = 0.04,  
n\_estimators = 100





## Our Best Results (RMSLE)

	<b>Feature Engineering + GSCV ElasticNet</b>	<b>Neural Net</b>	<b>Gradient Boosting Regression</b>	<b>F_Regression</b>	<b>Ensembles (Stacking CV Regressor)</b>	<b>One Drop Lasso Model</b>
Cross Validation	0.110	0.172	0.119	0.132	0.108	/
Dev_X	0.114	0.153	0.120	0.129	0.111	0.139
Kaggle Test (80% Train data)	0.178	0.261	0.181	0.167	0.181	
Kaggle Test (100% Train data)	0.160	0.252	0.185	0.159	0.156	